1 · QUICK DIAGNOSIS OF WEEK 1

We finished 5 of 30 videos: rework in scripting and change churn in motion are choking the line

Units started vs. finished by stage (bar = 35 units)

	Stage	Owner	Finished	% finished	h/unit · blocked	Note
	1. Brief & content definition	Content Authors	35/35	100%	1 h/unit · 0 blocked	—
#1	2. Script creation & review	Scriptwriters	20/35	57%	4 h/unit · 15 blocked	Misalignment with authors; heavy rework
	3. Visual asset design	Graphic Designers	15/20	75%	3 h/unit · 5 blocked	Waiting for final script
#2	4. Motion graphics & animation	Motion Designers	8/15	53%	5 h/unit · 7 blocked	Frequent changes from earlier stages
	5. Final video (Synthesia + sound)	Synthesia / Sound	5/8	62%	6 h/unit · 3 blocked	A/V sync; last-minute visual tweaks

WEEK 1 OUTPUT
5 / 30
17% of weekly target

NEEDED W2–W10
32.8 / week
to still deliver 300 by Week 10

STUCK IN PROGRESS
30 units
19 touch-hours per finished video

#1 Script creation & review (Stage 2)

- **Evidence:** 15 of 35 units (43%) stuck, the largest loss in the line. It also starves Stage 3, where 5 units are waiting for a final script.
- **Root cause:** no shared 'definition of ready' between Content Authors and Scriptwriters. Reviews happen one after another, so each misunderstanding means a full rewrite (4 h/unit).
- **So what:** at most 20 scripts a week can enter production, so 30 a week is impossible upstream.

#2 Motion graphics & animation (Stage 4)

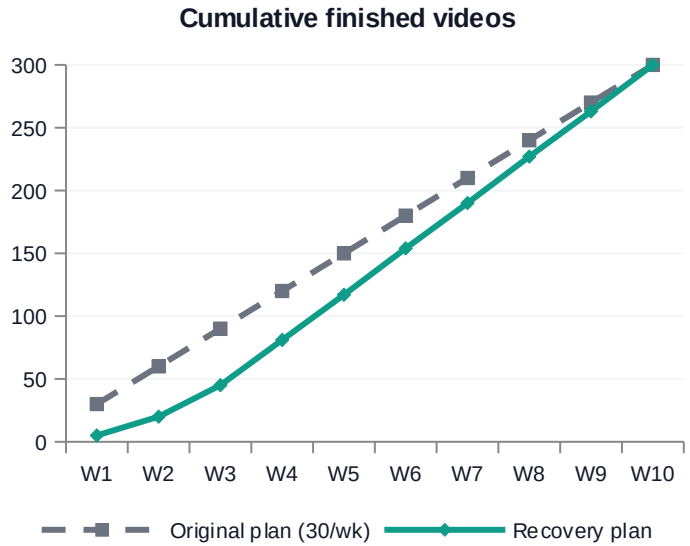
- **Evidence:** lowest completion rate (53%, 8 of 15) and the most effort of any creative stage (5 h/unit). 7 units are blocked by changes from earlier stages.
- **Root cause:** no freeze point. Script and visual edits keep arriving after animation has started, and every video is animated from scratch.
- **So what:** only 8 units reach final assembly. At target pace motion needs ~150 h/week; it lost a big part of Week 1 to redoing work.

Common thread: changes flow downstream unchecked, and Stage 1 keeps starting work the line can't absorb.

Luca · Chief of Staff hiring case · all data fictional | Source: outgoing PMO Week 1 tracking

Five rules for the next 2 weeks: align upstream, freeze changes, limit WIP, add motion capacity

#	Action (decision / rule we put in place)	Owner	When	Success metric (end of Week 3)
1	**One brief, co-signed.** The Content Author and Scriptwriter fill in a 1-page 'script-ready brief' in a 20-min pairing before any script starts: learning objective, key concepts, scene outline, 600–850 words. Author reviews within 24 h, max 2 rounds; a 3rd round goes to the Curriculum Lead to decide.	Head of Content Ops	Day 1–2	≥80% scripts approved first pass; ≤1.2 rework loops / unit
2	**Lock points + change control.** The script locks when it leaves Stage 2 and visuals lock when they leave Stage 4. After a lock, only factual or pedagogical errors may change (Content Lead approves within 24 h). Everything else goes to a v2 backlog.	Curriculum Lead	Day 1	0 change-driven blockers in Stages 4–5
3	**Stop starting, start finishing.** Cap the queue in front of each stage at 10 units. Stage 1 pauses new briefs while the script queue is over 10. Authors spend 50% of their time in pair-fix sessions to clear the 15 blocked scripts.	PMO / Chief of Staff	Day 1–5	Script backlog 15 → <5; WIP ≤ 40 units
4	**Add capacity where it's scarce.** 3 complexity tiers (A bespoke 20% · B template 50% · C avatar + static 30%). A library of 15–20 reusable motion templates. 2–3 freelance motion designers for W2–W6. Designers start from the approved scene outline, in parallel with the script.	Motion Lead + Head of Content Ops	Week 1–2	Motion 5.0 → ≤3.0 h/unit; ≥36 units/week out of Stage 4
5	**Right-first-time final assembly.** One technical spec (audio −16 LUFS, 48 kHz, timings from the locked script), a Synthesia template per subject, sound + Synthesia working as one cell on the same batch, and a 10-point QC checklist before export.	Synthesia Lead	Week 2	0 A/V sync rejects; Stage 5 6 → 4 h/unit



Weekly finish targets

W2: 15 · W3: 25 · W4–W10: ≥36/week → 300 by Week 10.

W2 mostly clears the 30 units already in progress; from W4 the tiers and templates carry the higher pace.

Escalation rule: if W3 is under 20, the steering group moves more videos to Tier C and orders subjects by the school-term calendar.

One AI-orchestrated pipeline from curriculum to finished video, with humans approving at key gates

KNOWLEDGE LAYER (RAG) SEP curriculum & learning objectives · Luca lesson library & pedagogy style guide · brand/asset & motion-template library · performance data (watch time, quiz results)

1. Brief agent	2. Script agent + QA	3. Scene spec (JSON)	4. Visual assets	5. Motion renderer	6. Avatar, voice & finish
LLM + RAG drafts objective, concepts, misconceptions, scene outline	Writes in Luca template; auto-checks length, reading level, accuracy, safety	Per scene: narration, on-screen text, visual prompt, template ID, duration	Reuse library first; generate brand-styled images; diagrams & math as code	Fills motion templates from scene JSON; bespoke only for Tier A	Synthesia API render, voice-over, music, loudness, captions, sync check
Author approves	**Scriptwriter edits & locks**	**Single source of truth**	**Designer curates**	**Motion designer polishes**	**Final QA sign-off**
Effort/unit: 1.0 → 0.3 h	**Effort/unit: 4.0 → 1.5 h**	**Effort/unit: auto**	**Effort/unit: 3.0 → 1.5 h**	**Effort/unit: 5.0 → 2.0 h**	**Effort/unit: 6.0 → 2.0 h**

Layer	Engine & sourcing (buy best-of-breed APIs, build only the thin orchestration layer)
LLMs: brief, script, QA	Claude for long-form pedagogy in Spanish & Portuguese; GPT / Gemini through a model router for structured JSON and second-opinion QA. pgvector RAG over the curriculum; 50 gold scripts as eval set.
Visual assets	Adobe Firefly / Ideogram / Flux (commercially safe image generation) with brand style presets; SVG + Manim for diagrams and math; Figma asset library.
Motion	Remotion (programmatic React video) or After Effects templates rendered via Nexrender / Plainly from the scene JSON; Runway / Veo for short B-roll clips.
Avatar, voice, sound	Synthesia API (current tool, Luca avatars); ElevenLabs for voice-over and LatAm dubbing; licensed music API; ffmpeg loudness (EBU R128); Whisper captions + script-vs-audio sync check.
Orchestration & tracking	n8n for the MVP → Temporal at scale; status written to the master tracker via API; review UI in Lovable / Retool with comments per scene; S3 storage with versioning.

Impact, guardrails & prototype

- **≈19 h → ~7 h of human effort per video** (−60%), so the same team can deliver 2–3× the output.
- **Guardrails:** human sign-off on brief and final; child-safety, bias and accuracy checks; an IP/licence log per asset; cost per video tracked.
- **Prototype:** no link in this version. I propose a 2-week MVP (Lovable front end + Claude API): topic → brief → script → scene JSON → Synthesia draft, piloted on Primaria Math.

Every Friday we know whether we're on track: one tracker, five fields, three meetings

Master tracker: 5 key fields beyond Unit ID & Stage

Field	What it captures	How it spots delays / drives accountability
1. Accountable owner	One named person for the unit's current stage	No shared ownership: every late unit has a name attached, and the stand-up goes owner by owner.
2. Stage due date (SLA)	Stage start + SLA (e.g. Script 2 days, Motion 2 days)	Auto-calculates 'days late'; red the moment a unit misses its SLA, not at week end.
3. Days in stage	Today − stage start date (aging)	Shows silent stalls before anyone flags a blocker; >2 days turns amber.
4. Status + blocker reason	On track / At risk / Blocked + a standard reason list (e.g. content misalignment, change after lock, A/V sync)	Blockers by reason × stage show the system's failure modes; blocked >24 h is escalated.
5. Rework loops	Number of times the unit was sent back to an earlier stage	Measures quality at hand-offs; ≥2 loops triggers a root-cause look in the weekly review.

Weekly scoreboard (auto, in the tracker): finished vs. 30 target, cumulative vs. plan, run-rate needed to hit 300, WIP and blockers per stage, average days in stage.

Weekly operating rhythm (3 recurring meetings)

Daily production stand-up

Mon–Fri · 15 min

Who: Stage leads (content, script, design, motion, Synthesia/sound) + PMO
Reviews: Walk the board right-to-left (closest to done first): units blocked >24 h, queues over the WIP limit, late units.
Decides: Who unblocks what today; capacity swaps between stages; escalations to the Curriculum Lead.

Weekly pipeline review

Monday · 45 min

Who: Chief of Staff (chair), Head of Content Ops, Curriculum Lead, all stage leads
Reviews: Throughput vs. plan, cycle time per stage, rework rate, top blocker reasons, change requests, next week's capacity.
Decides: Weekly finish commitment per stage; WIP limits; tier mix; approve/deny post-lock changes; freelance hours.

Bi-weekly steering

Every other Thursday · 30 min

Who: CEO / COO, Head of Product / Content, Chief of Staff
Reviews: Forecast to 300 vs. term start, budget vs. actual, AI tool MVP progress, risks.
Decides: Scope trade-offs (subject priority, Tier C share), budget for capacity, AI tool go/no-go.

Plus a Friday 1-page status note (not a meeting) to leadership: on track / at risk, top 3 risks, asks.